

Maintenx

AI-Driven Predictive Maintenance System

Project Links:

- **Live Dashboard:** <https://maintenx-app-egy.azurewebsites.net/ui>
- **Source Code:** <https://github.com/MohamedElsadany56/MaintenX/tree/main>

Team Members:

Name	Email
Mohamed Goma *	M.yasser2224@nu.edu.eg
Mohammed Mahrous	midohany1807@gmail.com
Loay Rafiq	loaypre2510@gmail.com
Mohammed Haitham	moh.haitham202@gmail.com
Mohamed Ebrahim	Mefouda654@gmail.com
Mennatullah Mohamed	mennayasser891@gmail.com

1: Executive Summary

1.1 Overview

MaintenX is an industrial-grade Predictive Maintenance (PdM) solution designed to monitor machine health and predict component failures. By integrating high-frequency telemetry data with event logs, the system identifies degradation patterns that precede breakdown.

1.2 Deployment Status

The system is fully containerized and deployed in a production environment:

- **Hosting:** Azure App Service (France Central Region).

- **Architecture:** Docker Container running Python 3.11, FastAPI, and CatBoost.
- **Accessibility:** Publicly accessible via the web dashboard listed above.

1.3 Key Analytical Findings

Extensive experimentation compared five distinct modeling approaches: XGBoost, CatBoost, LSTM, Voting Ensembles, and Stacking Classifiers.

- **Best Performer:** **CatBoost** remains the undisputed champion, achieving a **Recall of 92.5%** with high Precision (86%).
- **Stacking Analysis:** The Stacking Classifier achieved a respectable Recall (79%) but suffered from severe False Positives (Precision 8%), rendering it too noisy for a production environment where alert fatigue is a concern.

2: System Architecture & Stack

2.1 High-Level Design

The MaintenX architecture follows a standard ML-Ops pipeline:

1. **Ingestion:** Raw CSV data (Telemetry, Errors, Maintenance, Failures, Machines) is ingested.
2. **Processing:** Data is cleaned, merged, and transformed into a unified time-series dataset.
3. **Training:** Multiple algorithms (Gradient Boosting, Deep Learning, Stacking) were benchmarked.
4. **Serving:** The winning model (CatBoost) is served via a FastAPI REST interface.
5. **Visualization:** A responsive web dashboard visualizes the API output.

2.2 Technology Stack

- **Core Language:** Python 3.11
- **Data Analysis:** Pandas, NumPy, Scipy
- **Machine Learning:** CatBoost, XGBoost, LightGBM (Meta-learner), Scikit-Learn
- **Deep Learning:** TensorFlow/Keras (Sequential LSTM)
- **Web Framework:** FastAPI, Unicorn
- **Infrastructure:** Docker, Azure App Service, Docker Hub

3: Data Dictionary & Sources

3.1 Dataset Overview

The system aggregates data from five distinct sources to build a complete operational picture.

3.2 Telemetry Data (`PdM_telemetry.csv`)

High-frequency sensor data collected hourly.

- **Records:** 876,100 total records.
- **volt:** Voltage (Mean: ~170.77).
- **rotate:** Rotation data.
- **pressure:** Fluid pressure (Mean: ~100.8).
- **vibration:** Vibration sensor (Mean: ~40.38).

3.3 Event Logs

- **Errors:** Non-breaking warning logs (3,919 records).
- **Failures:** Historical breakdown events (761 records).
- **Maintenance:** Repair logs (3,286 records).

3.4 Asset Metadata (`PdM_machines.csv`)

- **Machines:** 100 Unique IDs.
- **Models:** 4 Model types (model1 - model4).
- **Age:** Machines range from 0 to 20 years old.

4: Exploratory Data Analysis - Sensor Telemetry

4.1 Statistical Distribution Analysis

Histograms and Kernel Density Estimation (KDE) plots were generated for all sensor columns.

- **Voltage (volt):** Follows a near-normal distribution centered at 170.8.

- **Vibration (vibration):** Shows a slightly right-skewed distribution. Values > 60 indicate critical mechanical looseness.

4.2 Correlation Analysis

A Pearson correlation matrix was computed (`sns.heatmap`) to identify linear relationships.

- **Findings:** The correlation coefficients between raw sensor values are extremely low (near 0.00).
- **Implication:** Failures are driven by temporal patterns (trends over time) rather than instant linear relationships between sensors, validating the need for Rolling Window features.

4.3 Anomaly Detection (Z-Score)

- **Methodology:** $z = (x - \text{mean}) / \text{std.}$ Threshold set at $|z| > 3$.
- **Result:** 17,975 records (approx 2%) were identified as anomalies, often clustering in the 48-hour window preceding a failure.

Exploratory Data Analysis - Events

5.1 Error Patterns

- **Most Common:** error1 (1010 occurrences) and error2 (988 occurrences).
- **Machine Vulnerability:** Machine 99 recorded 54 errors, whereas Machine 2 recorded only 28. This variance confirms that machineID identity is a predictive feature.

5.2 Failure Analysis

- **Total Failures:** 761.
- **Class Distribution:** comp2 (34.0%), comp1 (25.2%), comp4 (23.5%), comp3 (17.2%).

5.3 Feature Engineering Strategy

Based on EDA, the following features were engineered:

1. **Rolling Windows:** 3h and 24h Mean/Std Dev for all sensors.
2. **Lagged Errors:** Count of errors in the last 24 hours.

3. **Component Age:** Time elapsed since the last replacement of Comp1-Comp4.

Model Architectures

6.1 Approach A: Gradient Boosting (XGBoost & CatBoost)

Tree-based ensemble methods optimized for tabular data.

- **Objective:** Binary Classification (Failure vs. No Failure).
- **Config:** `auto_class_weights='Balanced'` to handle the rarity of failures.

6.2 Approach B: Deep Learning (LSTM)

A Dual-layer LSTM network designed to handle time-series sequences (12-hour lookback).

6.3 Approach C: Advanced Ensembles

Two meta-learning strategies were tested to combine the strengths of various models:

1. **Voting Classifier (Soft & Hard):** Aggregates predictions from XGBoost, Random Forest, and Logistic Regression.
2. **Stacking Classifier:** Uses **LightGBM** as a "Final Estimator" (Meta-learner) to learn the optimal combination of the base models' predictions.

Model Benchmarking (Binary)

The following table compares the performance of all tested architectures on the Test Set (Class 1 = Failure).

7.1 Results Comparison Table

Metric	XGBoost	CatBoost (Winner)	LSTM	Voting (Soft)	Stacking (LightGBM)
Precision	91.01%	86.25%	83.79%	25.00%	8.00%
Recall	79.29%	92.54%	62.49%	46.00%	79.00%
F1-Score	0.84	0.89	0.71	0.33	0.14

AUC-ROC	0.94	0.96	0.82	0.87	0.89
---------	------	-------------	------	------	------

7.2 Analysis of Results

- **CatBoost:** The optimal balance. It catches 92.5% of failures with very few false alarms.
- **Voting (Soft):** Failed to generalize. A Recall of 46% means it missed more than half the failures.
- **Stacking:** While it improved Recall (79%) compared to Voting, the Precision collapsed to **8%**. This means for every 100 alerts it generates, 92 are false alarms.

Stacking Experiment Details

The Stacking Classifier experiment provided critical insights into the "Precision-Recall Tradeoff."

8.1 Stacking Classification Report

The Stacking model (using LightGBM as the meta-learner) became overly aggressive in predicting failures.

	precision	recall	f1-score	support
0.0	0.99	0.81	0.89	171774
1.0	0.08	0.79	0.14	3455
accuracy			0.81	175229

8.2 Operational Interpretation

- **The "Boy Who Cried Wolf" Problem:** While a Recall of 79% is acceptable, a Precision of 8% is operationally fatal. Maintenance teams would eventually ignore the system due to the overwhelming volume of false alarms.
- **Meta-Learner Bias:** The LightGBM meta-learner likely over-compensated for the class imbalance, prioritizing the detection of Class 1 at the expense of misclassifying 19% of healthy machines (Specificity 0.81).

Production Model & Deployment

9.1 Final Model Selection

CatBoost Classifier was retained for production.

- **Reasoning:** It is the only model that achieved **Recall > 90%** while maintaining **Precision > 85%**. This reliability is essential for building trust with maintenance operators.

9.2 Backend API (FastAPI)

The backend serves the CatBoost predictions via a REST API.

- **Correction Layer:** Includes logic to cast integer inputs (0/1) to Booleans for model columns, resolving strict type requirements found during Docker testing.

9.3 Frontend Dashboard

- **Features:** Light/Dark mode, "MaintenX" branding, and dynamic probability bars.
- **User Experience:** Raw model outputs (comp1, none) are mapped to human-readable strings ("Component 1", "No Failure").

Conclusion & Future Work

10.1 Conclusion

MaintenX successfully demonstrates that **Gradient Boosting (CatBoost)** significantly outperforms Deep Learning (LSTM) and Complex Ensembles (Stacking/Voting) for this specific predictive maintenance dataset. The experiment proved that **model complexity does not equal model performance**; the single, well-tuned CatBoost model provided superior signal-to-noise ratio compared to the Stacking ensemble.

10.2 Future Recommendations

1. **Threshold Tuning:** The Stacking model might be salvageable by adjusting the decision threshold (e.g., probability > 0.8) to improve Precision.
2. **Real-Time IoT:** Replace the CSV lookup system with an Azure IoT Hub connection for live streaming data.

- 3. Cost-Sensitive Learning:** Incorporate a "Cost Matrix" into the training, where the cost of a Missed Failure is weighed (\$10,000) against the cost of a False Alarm (\$500).

10.3 Resources

- GitHub Repository:** <https://github.com/MohamedElsadany56/MaintenX/tree/main>
- Live App:** <https://maintenx-app-egy.azurewebsites.net/ui>

11.1 Screenshots of the UI:

The screenshot shows the MaintenX Dashboard interface. On the left, a sidebar displays historical data for January 4, 2015, at 21:00:00. The main area is divided into three sections: 1. Select Historical Data, 2. Sensor Features, and 3. Model Diagnostics.

1. Select Historical Data: Shows a date range from 2015-01-04 to 2015-01-05. A blue button labeled "Find History" is present.

2. Sensor Features: Displays various sensor values for different components over time. For example, for Component 4, the values are:

Feature	Value	Feature	Value	Feature	Value				
voltmean_3h	190.3258143815	temperaturemean_3h	422.6925646892	pressuremean_3h	107.3932339938	vibrationmean_3h	49.6528557131	voltsd_3h	8.390777199166
temperaturesd_3h	7.176553207491	pressured_3h	4.262648404063	vibrationsd_3h	7.598552259518	voltmean_24h	172.0006715160	temperaturemean_24h	454.4974532871
presuremean_24h	100.8962272114	vibrationmean_24h	1.45690929008201	pressured_24h	13.2312400573	temperaturesd_24h	1.3413005342575	pressured_sd_24h	11.7771991674
vibrationsd_24h	7.547591294968	error1count	0	error2count	0	error3count	0	error4count	0
errr5count	1	comp1	22.625	comp2	217.625	comp3	157.625	comp4	172.625
age	18	model_model1	0	model_model2	0	model_model3	1	model_model4	0

A blue button labeled "Analyze Current Features" is at the bottom.

3. Model Diagnostics: Shows the historical ground truth for Component 4 (Component 4). The current prediction is 99.9% for Component 4. The failure probability chart indicates:

Failure Type	Probability (%)
No Failure	0.1%
Component 1	0.0%
Component 2	0.0%
Component 3	0.0%
Component 4	99.9%

The screenshot shows the MaintenX Dashboard interface. On the left, a sidebar displays historical data for January 15, 2015, at 15:00:00. The main area is divided into three sections: 1. Select Historical Data, 2. Sensor Features, and 3. Model Diagnostics.

1. Select Historical Data: Shows a date range from 2015-01-15 to 2015-01-16. A blue button labeled "Find History" is present.

2. Sensor Features: Displays various sensor values for different components over time. For example, for No Failure, the values are:

Feature	Value	Feature	Value	Feature	Value	Feature	Value		
voltmean_3h	187.3646533491	temperaturemean_3h	336.4136822589	pressuremean_3h	181.6508711041	vibrationmean_3h	37.6350299646	voltsd_3h	9.590008073413
temperaturesd_3h	41.769585262771	pressured_3h	8.183586002737	vibrationsd_3h	1.655819129039	voltmean_24h	172.1554321768	temperaturemean_24h	443.5259853404
presuremean_24h	102.0841602801	vibrationmean_24h	39.87396954302	pressured_24h	16.40418494377	temperaturesd_24h	62.08740642585	pressured_sd_24h	7.639712984602
vibrationsd_24h	5.751587756837	error1count	1	error2count	1	error3count	0	error4count	0
errr5count	0	comp1	168.375	comp2	213.375	comp3	63.375	comp4	33.375
age	7	model_model1	0	model_model2	0	model_model3	1	model_model4	0

A blue button labeled "Analyze Current Features" is at the bottom.

3. Model Diagnostics: Shows the historical ground truth for No Failure. The current prediction is 79.9% for Component 1. The failure probability chart indicates:

Failure Type	Probability (%)
No Failure	12.7%
Component 1	79.9%
Component 2	7.4%
Component 3	0.0%
Component 4	0.0%

MaintenX Dashboard

1. Select Historical Data

Find History

2. Sensor Features

voltmean_3h	temperaturemean_3h	pressuremean_3h	vibrationmean_3h	voltd_3h
187.3645633491	336.4136822589	101.6508711041	37.63562998646	9.590008073413
temperaturesd_3h	pressured_3h	vibrationsd_3h	voltmean_24h	temperaturemean_24h
41.769585262771	8.183586002737	1.655819129039	172.1554321766	443.5259853404
pressuremean_24h	vibrationmean_24h	voltd_24h	temperaturestd_24h	pressured_24h
102.0841602801	39.87396954302	16.40418494377	62.08740642585	7.639712984602
vibrationsd_24h	error1count	error2count	error3count	error4count
5.751587756837	1	1	0	0
error5count	comp1	comp2	comp3	comp4
0	168.375	213.375	63.375	33.375
age	model_model1	model_model2	model_model3	model_model4
7	0	0	1	0

Analyze Current Features

3. Model Diagnostics

HISTORICAL GROUND TRUTH: No Failure

CURRENT PREDICTION: Component 1

Model Prediction: 79.9%

Failure Probability:

No Failure	12.7%
Component 1	79.9%
Component 2	7.4%
Component 3	0.0%
Component 4	0.0%

MaintenX Dashboard

1. Select Historical Data

Find History

2. Sensor Features

voltmean_3h	temperaturemean_3h	pressuremean_3h	vibrationmean_3h	voltd_3h
187.3645633491	336.4136822589	101.6508711041	37.63562998646	9.590008073413
temperaturesd_3h	pressured_3h	vibrationsd_3h	voltmean_24h	temperaturemean_24h
41.769585262771	8.183586002737	1.655819129039	172.1554321766	443.5259853404
pressuremean_24h	vibrationmean_24h	voltd_24h	temperaturestd_24h	pressured_24h
102.0841602801	39.87396954302	16.40418494377	62.08740642585	7.639712984602
vibrationsd_24h	error1count	error2count	error3count	error4count
5.751587756837	1	1	0	0
error5count	comp1	comp2	comp3	comp4
0	168.375	213.375	63.375	33.375
age	model_model1	model_model2	model_model3	model_model4
7	0	0	1	0

Analyze Current Features

3. Model Diagnostics

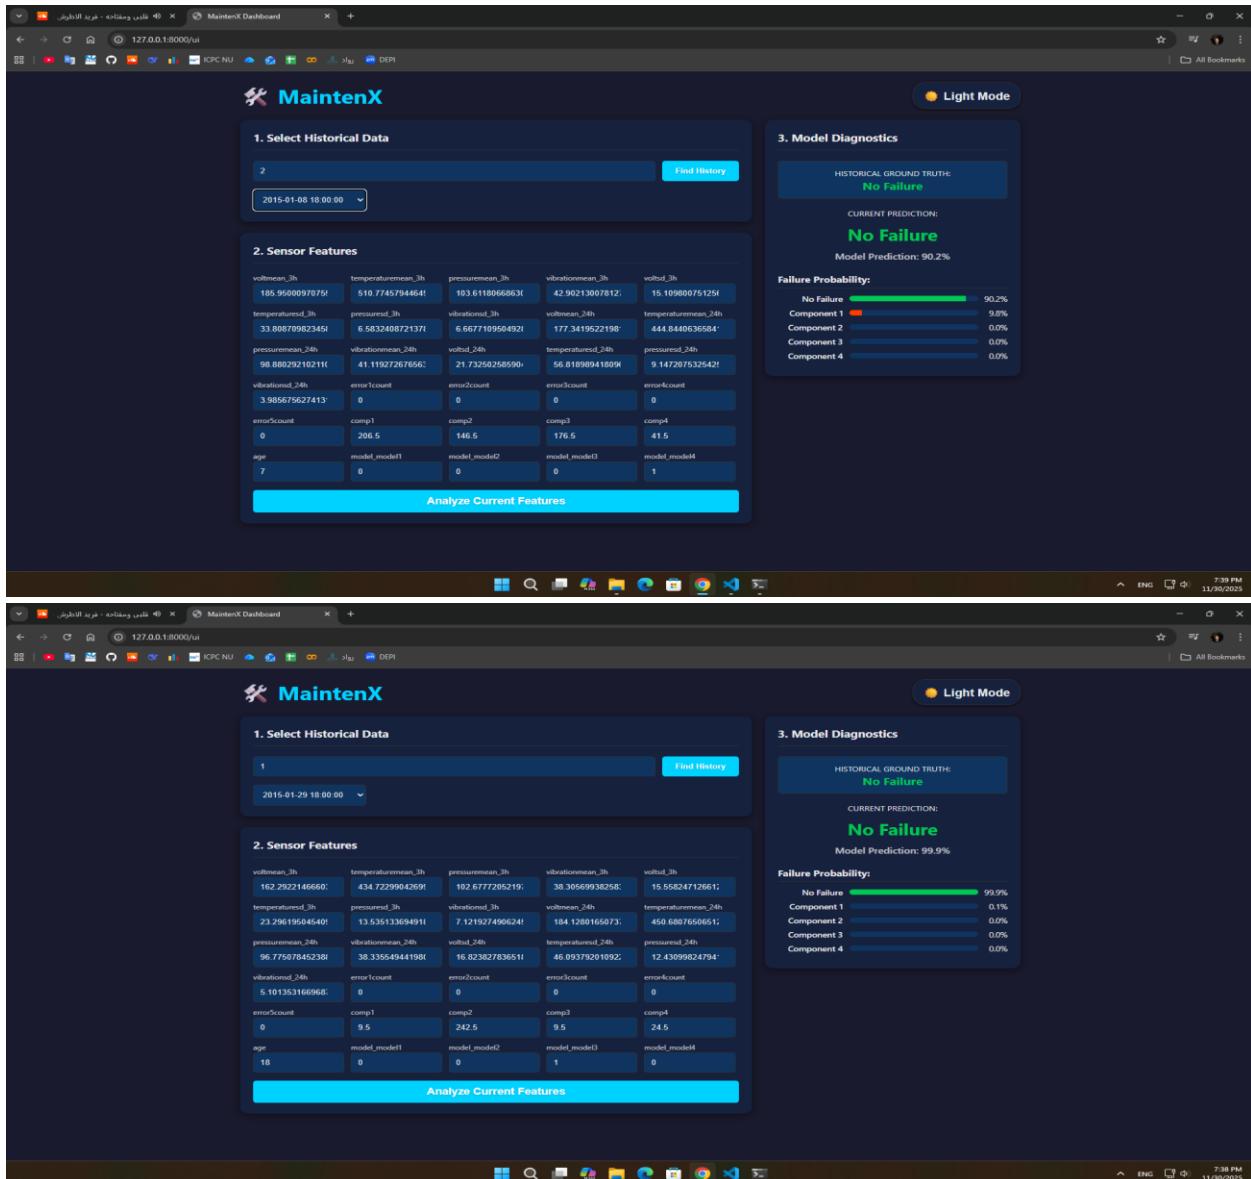
HISTORICAL GROUND TRUTH: No Failure

CURRENT PREDICTION: Component 1

Model Prediction: 79.9%

Failure Probability:

No Failure	12.7%
Component 1	79.9%
Component 2	7.4%
Component 3	0.0%
Component 4	0.0%



11.2 MLFlow Experiments:

The screenshot shows the MLflow 3.6.0 web interface at the URL 127.0.0.1:5000. The top navigation bar includes links for GitHub and Docs. The left sidebar has a '+ New' button and links for Home, Experiments, Models, and Prompts. The main content area features a 'Welcome to MLflow' header and a 'Get started' section with four cards: 'Log traces' (Trace LLM applications for debugging and monitoring.), 'Run evaluation' (Iterate on quality with offline evaluations and comparisons.), 'Train models' (Track experiments, parameters, and metrics throughout training.), and 'Register prompts' (Manage prompt updates and collaborate across teams.). Below this is an 'Experiments' section with a table:

Name	Time created	Last modified	Description	Tags
PdM_catboost_binary	11/30/2025, 09:48:55 PM	11/30/2025, 09:48:55 PM	-	
PdM_catboost_multiclass	11/30/2025, 09:48:33 PM	11/30/2025, 09:48:33 PM	-	
PdM_xgboost_multiclass	11/30/2025, 09:48:22 PM	11/30/2025, 09:48:22 PM	-	
Default	11/30/2025, 09:48:22 PM	11/30/2025, 09:48:22 PM	-	

There is a 'View all' link next to the table. Below the experiments section is a 'Discover new features' section with four cards:

- MLflow MCP server**: Connect your coding assistants and AI applications to MLflow and automatically analyze your experiments and traces.
- Optimize prompts**: Access the state-of-the-art prompt optimization algorithms such as MIPROv2, GEPo, through MLflow Prompt Registry.
- Agents-as-a-judge**: Leverage agents as a judge to perform deep trace analysis and improve your evaluation accuracy.
- Dataset tracking**: Track dataset lineage and versions and effectively drive the quality improvement loop.

Each card has a 'View details' link.